

Project - Loan Default

Table of Contents

- [Description of Project](#)
- [Dataset Loading and Analysis](#)
 - [Data Loading](#)
 - [Data Cleaning](#)
 - [Dummy Variables](#)
 - [Data Imbalance - SMOTE](#)
 - [Cross Validation](#)
- [Model Fitting and Evaluation](#) with Confusion Matrix, RoC Curve (& Precision Recall), Calibration Plot, Type of Errors
 - [Logit Regression](#)
 - [kNN](#)
 - [Random Forest](#)
 - [XGBoost](#)
 - [Compare ROC curve](#)
 - [Compare Precision-Recall curve](#)

Description

- General overview of project: background and what we want to solve
- Dataset used (source):
https://machinehack.com/hackathons/deloitte_presents_machine_learning_challeng

Dataset Loading and Analysis

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from imblearn.over_sampling import SMOTE
import imblearn
from sklearn.datasets import make_classification, load_iris
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as L
from sklearn.dummy import DummyClassifier
from sklearn.metrics import roc_curve, auc, confusion_matrix, precision_r
from jupyterthemes import jtplot
from sklearn.preprocessing import StandardScaler, MinMaxScaler, FunctionT
```

```
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.calibration import CalibratedClassifierCV, calibration_curve
from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

random_state=809
%matplotlib inline
plt.style.use('seaborn-v0_8-whitegrid') # If error reported by this line,
#jupyterplot.style(theme='onedork')
```

Data Loading

Since the initial dataset has already been seperated into train.csv and test.csv, we just load the data and don't need to seperate the in-sample data and out-of-sample data from train.csv. Here, train.csv data is regarded as in-sample data, while test.csv data is regarded as out-of-sample data

```
In [2]: df= pd.read_csv("./train.csv")

In [3]: df.head(10)
```

Out [3]:

| | ID | Loan Amount | Funded Amount | Funded Amount Investor | Term | Batch Enrolled | Interest Rate | Gra |
|---|----------|-------------|---------------|------------------------|------|----------------|---------------|-----|
| 0 | 65087372 | 10000 | 32236 | 12329.362860 | 59 | BAT2522922 | 11.135007 | |
| 1 | 1450153 | 3609 | 11940 | 12191.996920 | 59 | BAT1586599 | 12.237563 | |
| 2 | 1969101 | 28276 | 9311 | 21603.224550 | 59 | BAT2136391 | 12.545884 | |
| 3 | 6651430 | 11170 | 6954 | 17877.155850 | 59 | BAT2428731 | 16.731201 | |
| 4 | 14354669 | 16890 | 13226 | 13539.926670 | 59 | BAT5341619 | 15.008300 | |
| 5 | 50509046 | 34631 | 30203 | 8635.931613 | 36 | BAT4694572 | 17.246986 | |
| 6 | 32737431 | 30844 | 19773 | 15777.511830 | 59 | BAT4808022 | 10.731432 | |
| 7 | 63151650 | 20744 | 10609 | 7645.014802 | 58 | BAT2558388 | 13.993688 | |
| 8 | 4279662 | 9299 | 11238 | 13429.456610 | 59 | BAT5341619 | 11.178457 | |
| 9 | 4431034 | 19232 | 8962 | 7004.097481 | 58 | BAT2078974 | 5.520413 | |

10 rows x 35 columns

Data Cleaning

Clear NaN or NA Values (if any)

```
In [4]: # Find the count and percentage of missing values
df_na = pd.DataFrame({'Percent': 100*df.isnull().sum()/len(df), 'Count':

# Print columns with null count > 0
df_na[df_na['Count'] > 0]
```

Out[4]:

| Percent | Count |
|---------|-------|
|---------|-------|

```
In [5]: hasNA = False
for i in df.isna().sum():
    if i != 0:
        hasNA = True
print(hasNA)
```

False

There's no NA/NaN value in this dataset.

```
In [6]: df.dtypes
```

```
Out[6]: ID                                int64
Loan Amount                             int64
Funded Amount                           int64
Funded Amount Investor                   float64
Term                                    int64
Batch Enrolled                           object
Interest Rate                           float64
Grade                                    object
Sub Grade                               object
Home Ownership                           object
Salary                                  float64
Verification Status                       object
Payment Plan                             object
Loan Title                               object
Debit to Income                           float64
Delinquency - two years                    int64
Inquires - six months                     int64
Open Account                             int64
Public Record                             int64
Revolving Balance                         int64
Revolving Utilities                       float64
Total Accounts                           int64
Initial List Status                       object
Total Received Interest                   float64
Total Received Late Fee                   float64
Recoveries                               float64
Collection Recovery Fee                   float64
Collection 12 months Medical              int64
Application Type                          object
Last week Pay                             int64
Accounts Delinquent                       int64
Total Collection Amount                   int64
Total Current Balance                     int64
Total Revolving Credit Limit              int64
Loan Status                               int64
dtype: object
```

First, drop some obvious irrelevant columns. "ID" and "Batch Enrolled" are just identification number of loan, not useful for prediction.


```
In [7]: df.drop(['ID', 'Batch Enrolled'],axis=1,inplace=True)
```

Data Distribution

Correlation Analysis

```
In [8]: corr = df.corr(numeric_only=True )
# Plot the heatmap
plt.figure(figsize=(14,14))
sns.heatmap(corr,
            xticklabels=corr.columns,
            yticklabels=corr.columns,
            cmap='coolwarm')
```

Out[8]: <AxesSubplot: >



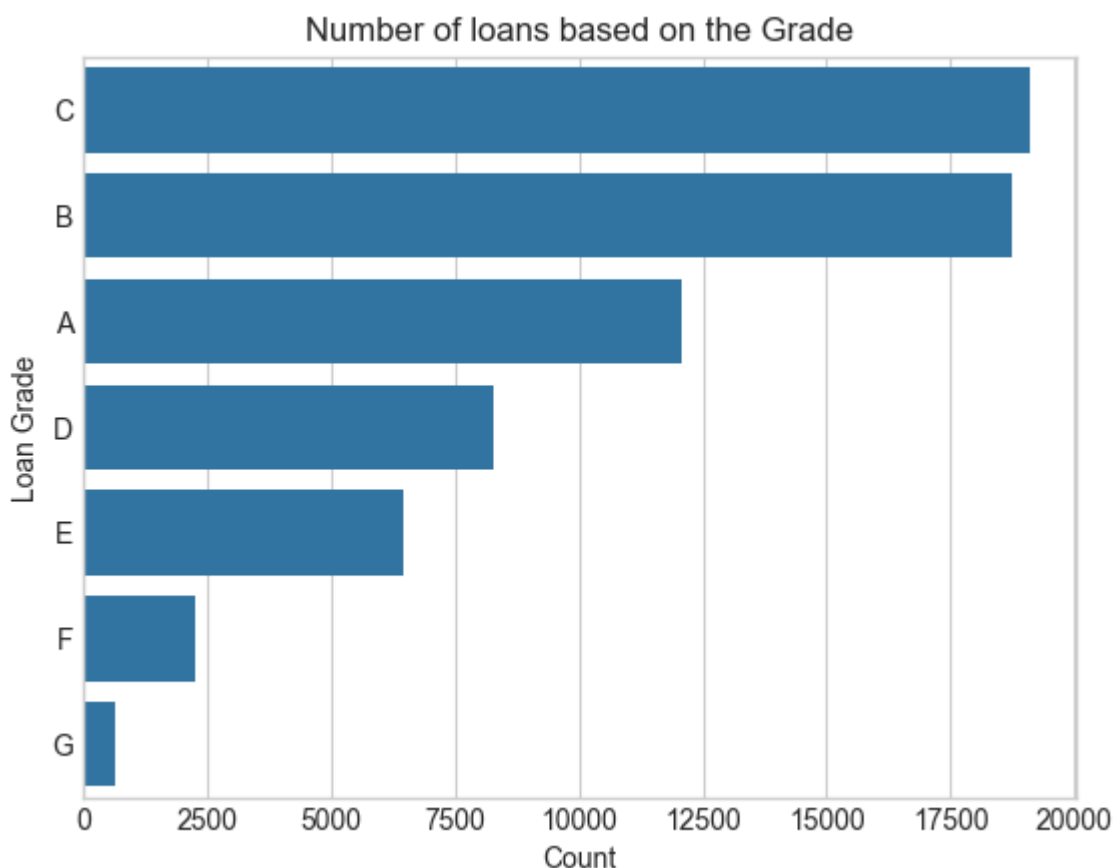
Based on the heatmap, the numerical values in the dataset does not reveal strong correlation.

Loan Grade

```
In [9]: grade_vis = df['Grade'].value_counts()

sns.barplot(y=grade_vis.index, x=grade_vis)
plt.title('Number of loans based on the Grade')
plt.ylabel('Loan Grade')
plt.xlabel('Count')
```

Out[9]: Text(0.5, 0, 'Count')



Drop columns

```
In [10]: df.drop(['Term', 'Verification Status', 'Payment Plan', 'Loan Title', 'Delinq
```

```
In [11]: df.drop(['Open Account', 'Total Accounts', 'Total Received Late Fee', 'Recov
```

```
In [12]: df.columns
```

```
Out[12]: Index(['Loan Amount', 'Funded Amount', 'Funded Amount Investor',
               'Interest Rate', 'Grade', 'Sub Grade', 'Home Ownership', 'Salar
y',
               'Debit to Income', 'Revolving Balance', 'Total Received Interes
t',
               'Total Collection Amount', 'Total Current Balance',
               'Total Revolving Credit Limit', 'Loan Status'],
              dtype='object')
```

Dummy Variables

```
In [13]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 67463 entries, 0 to 67462
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Loan Amount                          67463 non-null  int64
1   Funded Amount                        67463 non-null  int64
2   Funded Amount Investor               67463 non-null  float64
3   Interest Rate                        67463 non-null  float64
4   Grade                                67463 non-null  object
5   Sub Grade                            67463 non-null  object
6   Home Ownership                       67463 non-null  object
7   Salary                               67463 non-null  float64
8   Debit to Income                      67463 non-null  float64
9   Revolving Balance                    67463 non-null  int64
10  Total Received Interest               67463 non-null  float64
11  Total Collection Amount               67463 non-null  int64
12  Total Current Balance                 67463 non-null  int64
13  Total Revolving Credit Limit          67463 non-null  int64
14  Loan Status                           67463 non-null  int64
dtypes: float64(5), int64(7), object(3)
memory usage: 7.7+ MB

In [14]: categorical_variables=['Grade','Sub Grade',"Home Ownership"]
df_categorical=pd.get_dummies(data=df, columns=categorical_variables,pref

In [15]: df_categorical.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 67463 entries, 0 to 67462
```

```
Data columns (total 54 columns):
```

| # | Column | Non-Null Count | Dtype |
|----|------------------------------|----------------|---------|
| 0 | Loan Amount | 67463 non-null | int64 |
| 1 | Funded Amount | 67463 non-null | int64 |
| 2 | Funded Amount Investor | 67463 non-null | float64 |
| 3 | Interest Rate | 67463 non-null | float64 |
| 4 | Salary | 67463 non-null | float64 |
| 5 | Debit to Income | 67463 non-null | float64 |
| 6 | Revolving Balance | 67463 non-null | int64 |
| 7 | Total Received Interest | 67463 non-null | float64 |
| 8 | Total Collection Amount | 67463 non-null | int64 |
| 9 | Total Current Balance | 67463 non-null | int64 |
| 10 | Total Revolving Credit Limit | 67463 non-null | int64 |
| 11 | Loan Status | 67463 non-null | int64 |
| 12 | Col_B | 67463 non-null | int8 |
| 13 | Col_C | 67463 non-null | int8 |
| 14 | Col_D | 67463 non-null | int8 |
| 15 | Col_E | 67463 non-null | int8 |
| 16 | Col_F | 67463 non-null | int8 |
| 17 | Col_G | 67463 non-null | int8 |
| 18 | Col_A2 | 67463 non-null | int8 |
| 19 | Col_A3 | 67463 non-null | int8 |
| 20 | Col_A4 | 67463 non-null | int8 |
| 21 | Col_A5 | 67463 non-null | int8 |
| 22 | Col_B1 | 67463 non-null | int8 |
| 23 | Col_B2 | 67463 non-null | int8 |
| 24 | Col_B3 | 67463 non-null | int8 |
| 25 | Col_B4 | 67463 non-null | int8 |
| 26 | Col_B5 | 67463 non-null | int8 |
| 27 | Col_C1 | 67463 non-null | int8 |
| 28 | Col_C2 | 67463 non-null | int8 |
| 29 | Col_C3 | 67463 non-null | int8 |
| 30 | Col_C4 | 67463 non-null | int8 |
| 31 | Col_C5 | 67463 non-null | int8 |
| 32 | Col_D1 | 67463 non-null | int8 |
| 33 | Col_D2 | 67463 non-null | int8 |
| 34 | Col_D3 | 67463 non-null | int8 |
| 35 | Col_D4 | 67463 non-null | int8 |
| 36 | Col_D5 | 67463 non-null | int8 |
| 37 | Col_E1 | 67463 non-null | int8 |
| 38 | Col_E2 | 67463 non-null | int8 |
| 39 | Col_E3 | 67463 non-null | int8 |
| 40 | Col_E4 | 67463 non-null | int8 |
| 41 | Col_E5 | 67463 non-null | int8 |
| 42 | Col_F1 | 67463 non-null | int8 |
| 43 | Col_F2 | 67463 non-null | int8 |
| 44 | Col_F3 | 67463 non-null | int8 |
| 45 | Col_F4 | 67463 non-null | int8 |
| 46 | Col_F5 | 67463 non-null | int8 |
| 47 | Col_G1 | 67463 non-null | int8 |
| 48 | Col_G2 | 67463 non-null | int8 |
| 49 | Col_G3 | 67463 non-null | int8 |
| 50 | Col_G4 | 67463 non-null | int8 |
| 51 | Col_G5 | 67463 non-null | int8 |
| 52 | Col_OWN | 67463 non-null | int8 |
| 53 | Col_RENT | 67463 non-null | int8 |

dtypes: float64(5), int64(7), int8(42)
memory usage: 8.9 MB


Data Imbalance - SMOTE

```
In [16]: loan_status_val = df['Loan Status'].value_counts()
sns.barplot(y=loan_status_val, x=loan_status_val.index)
plt.title('Number of loans based on the Loan Status')
plt.ylabel('Count')
plt.xlabel('Loan Status')

print("Number of 0 observations (Loan Non-Defaulted):", loan_status_val[0])
print("Number of 1 observations (Loan Defaulted):", loan_status_val[1])
```

Number of 0 observations (Loan Non-Defaulted): 61222


Number of 1 observations (Loan Defaulted): 6241



```
In [17]: X = df_categorical.drop('Loan Status', axis=1)
y = df_categorical['Loan Status']
```

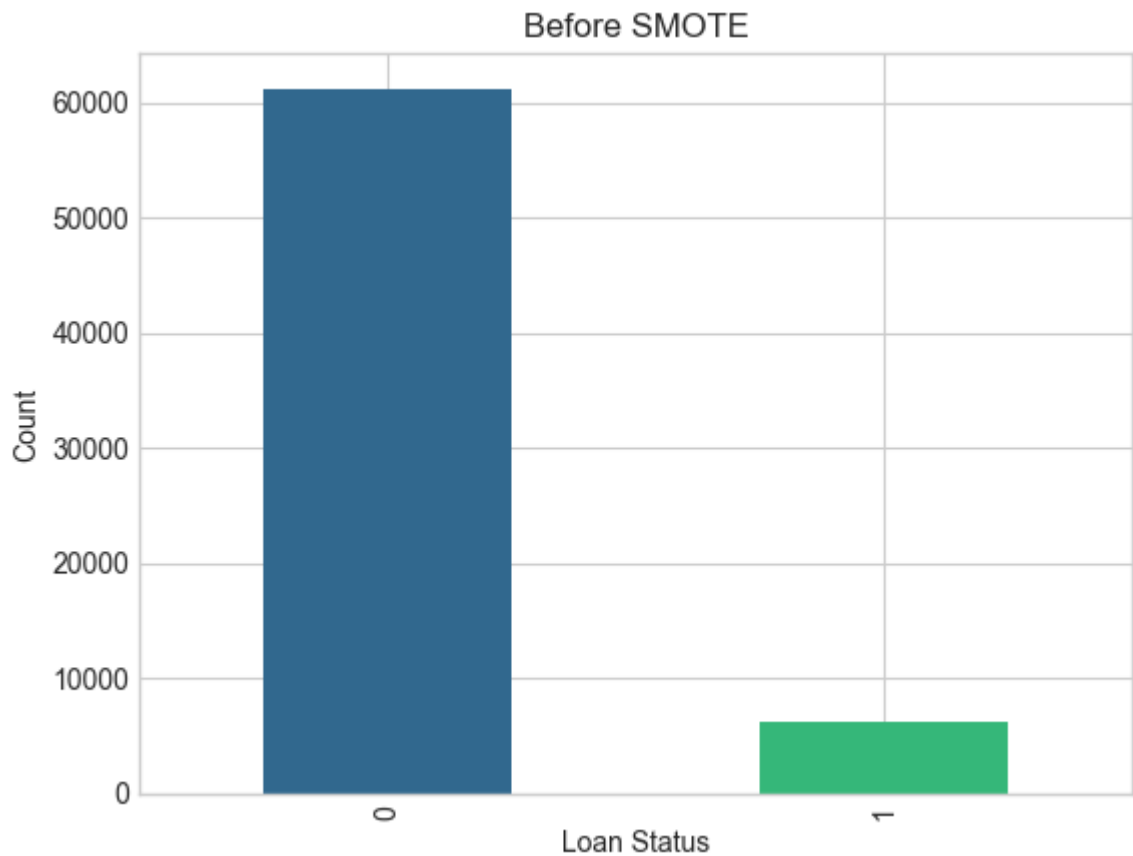
```
In [18]: # Visualize the dataset before SMOTE
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
sns.scatterplot(x=X.iloc[:, 0], y=X.iloc[:, 1], hue=y, palette='viridis',
plt.title('Dataset Before SMOTE')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
```

```
Out[18]: Text(0, 0.5, 'Feature 2')
```

```
In [23]: # Use matplotlib directly to control colors, if needed
colors = sns.color_palette('viridis', len(y.unique())) # Assuming y has
y.value_counts().plot(kind='bar', color=colors)
plt.title('Before SMOTE')
plt.xlabel('Loan Status')
plt.ylabel('Count')
```

```
Out[23]: Text(0, 0.5, 'Count')
```




```
In [24]: # Apply SMOTE
smote = SMOTE(random_state=random_state)
smote.fit(X,y)
X,y=smote.fit_resample(X,y)
```

```
In [25]: # Feature Scaling
scaler = MinMaxScaler()
x_scaled = scaler.fit_transform(X)
```

```
In [26]: # Visualize the dataset after SMOTE
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 2)
sns.scatterplot(x=X.iloc[:, 0], y=X.iloc[:, 1], hue=y, palette='viridis',
plt.title('Dataset After SMOTE')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
```

```
Out[26]: Text(0, 0.5, 'Feature 2')
```



```
In [27]: # Bar chart for class distribution after SMOTE
plt.figure(figsize=(8, 5))
plt.bar(['Class 0', 'Class 1'], [sum(y == 0), sum(y == 1)], color=['skybl
plt.title('Class Distribution After SMOTE')
plt.xlabel('Class')
plt.ylabel('Count')

# Show the plots
plt.show()
```



```
In [28]: # Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(x_scaled, y, test_size=0.2)
```

Cross validation

```
In [29]: # visualize the training and testing splits generated by a cross-validation
def plot_cv_indices(cv, X, y, n_splits, lw=10):
    fig, ax = plt.subplots(figsize = (15,8))
    # Generate the training/testing visualizations for each CV split
    for ii, (tr, tt) in enumerate(cv.split(X=X, y=y)):
        # Fill in indices with the training/test groups
        indices = np.array([np.nan] * len(X))
        indices[tt] = 1
        indices[tr] = 0

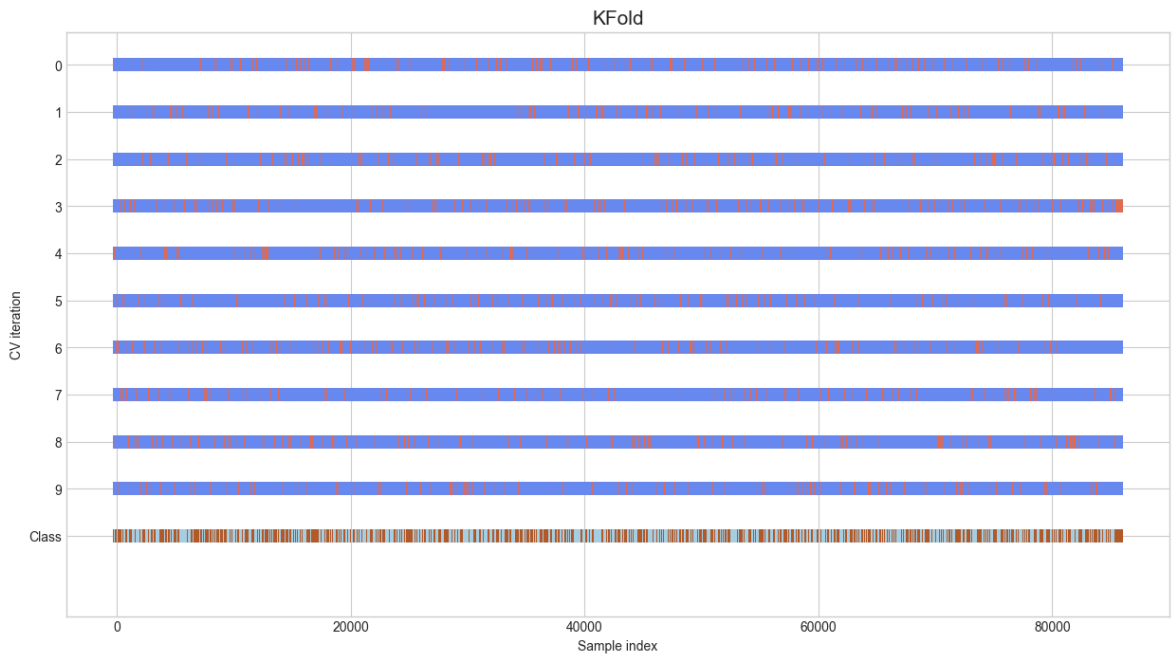
        # Visualize the results
        ax.scatter(range(len(indices)), [ii + .5] * len(indices),
                    c=indices, marker='_', lw=lw, cmap=plt.cm.coolwarm,
                    vmin=-.2, vmax=1.2)

    # Plot the data classes
    ax.scatter(range(len(X)), [ii + 1.5] * len(X), c=y, marker='_', lw=lw)

    # Formatting
    yticklabels = list(range(n_splits)) + ['Class']
    ax.set(yticks=np.arange(n_splits+1) + .5, yticklabels=yticklabels,
           xlabel='Sample index', ylabel="CV iteration",
           ylim=[n_splits+2.2, -.2])
    ax.set_title('{}'.format(type(cv).__name__), fontsize=15)
    return ax
```

```
In [30]: n_splits = 10
shuffle = True
cv = KFold(n_splits=n_splits, shuffle=shuffle, random_state=random_state)
```

```
plot = plot_cv_indices(cv, X_train, y_train, n_splits)
plt.show()
```



Model Fitting and Evaluation


```
In [31]: # Compare model evaluation at different values of K by adjusting the value
k_range = range(1,20)
scores = []

for k in k_range:
    knn = KNeighborsClassifier(n_neighbors =k)
    knn.fit(X_train,y_train)
    scores.append(knn.score(X_test,y_test))

max_score = max(scores)
max_k = k_range[scores.index(max_score)]
```

```
In [32]: plt.figure(figsize = (8,6))
plt.xlabel('k')
plt.ylabel('accuracy')
plt.scatter(k_range, scores)
plt.xticks([0,5,10,15,20])
```

```
Out[32]: ([<matplotlib.axis.XTick at 0x2c825dd3f50>,
<matplotlib.axis.XTick at 0x2c825db2850>,
<matplotlib.axis.XTick at 0x2c825dd2c90>,
<matplotlib.axis.XTick at 0x2c825e0b090>,
<matplotlib.axis.XTick at 0x2c825e19590>],
[Text(0, 0, '0'),
Text(5, 0, '5'),
Text(10, 0, '10'),
Text(15, 0, '15'),
Text(20, 0, '20')])
```



```
In [33]: # Set the model list
models = []
models.append(('LR', LogisticRegression(random_state=random_state)))
models.append(('KNN', KNeighborsClassifier(n_neighbors=max_k)))
# models.append(('DTC', DecisionTreeClassifier(random_state=random_state)))
models.append(('RF', RandomForestClassifier(n_estimators=100, random_state=random_state)))
models.append(('XGB', XGBClassifier(use_label_encoder=False, eval_metric='logloss')))
```

Do Cross Validation using different models to do the comparison


```
In [34]: # Do cross validation using different models to do the comparison
n_splits = 10
scoring = 'accuracy'
results = []
tags = []
for tag, model in models:
    cv = KFold(n_splits=n_splits, random_state=random_state, shuffle=True)
    cv_results = cross_val_score(model, x_scaled, y, cv=cv, scoring=scoring)
    results += [cv_results]
    tags += [tag]
    print( tag + ": " + f'{cv_results.mean():.4f}' + ' (' + f'{cv_results
```

```
LR: 0.9213 (0.0020)
KNN: 0.9246 (0.0027)
RF: 0.9334 (0.0025)
XGB: 0.9243 (0.0019)
```

```
In [35]: # Plot the comparison result
fig_cv_compare = plt.figure()
fig_cv_compare.suptitle('Cross Validation Comparison')
ax = fig_cv_compare.add_subplot(111)
```



```
plt.boxplot(results)
ax.set_xticklabels(tags)
plt.show()
```



Do Model Accuracy using different models to do the comparison

```
In [36]: # Dictionary to store accuracy of each model
model_accuracies = {}
acc = []
tags = []
# Iterate over each model, train it, predict on the testing set, and calc
for tag, model in models:
    # Fit the model on the training set
    model.fit(X_train, y_train)

    # Predict on the testing set
    y_pred = model.predict(X_test)

    # Calculate the accuracy and store it
    accuracy = accuracy_score(y_test, y_pred)
    model_accuracies[tag] = accuracy
    acc += [accuracy]
    tags += [tag]
    print(f'{tag} Accuracy: {accuracy:.4f}')
```

```
LR Accuracy: 0.9206
KNN Accuracy: 0.9236
RF Accuracy: 0.9294
XGB Accuracy: 0.9233
```

```
In [37]: # Plot the comparison result as a bar chart
fig_acc_compare = plt.figure(figsize=(10, 6))
fig_acc_compare.suptitle('Accuracy Comparison of Different Models')
ax = fig_acc_compare.add_subplot(111)


# Create a bar chart
ax.bar(tags, acc, color=['blue', 'green', 'red', 'purple', 'orange'])

ax.set_ylabel('Accuracy')
ax.set_xlabel('Model')
ax.set_ylim(min(acc) - 0.05, max(acc) + 0.05) # Adjust y-axis limits to

# Add accuracy values on top of the bars
for i, accuracy in enumerate(acc):
    ax.text(i, accuracy + 0.01, f'{accuracy:.4f}', ha='center')

plt.show()
```

Accuracy Comparison of Different Models



Logistic Regression

Preparation

```
In [38]: # Initialize and Fit Logistic Regression Model
logit = LogisticRegression(random_state=random_state)

#Fit the model
logit.fit(X_train, y_train)
```

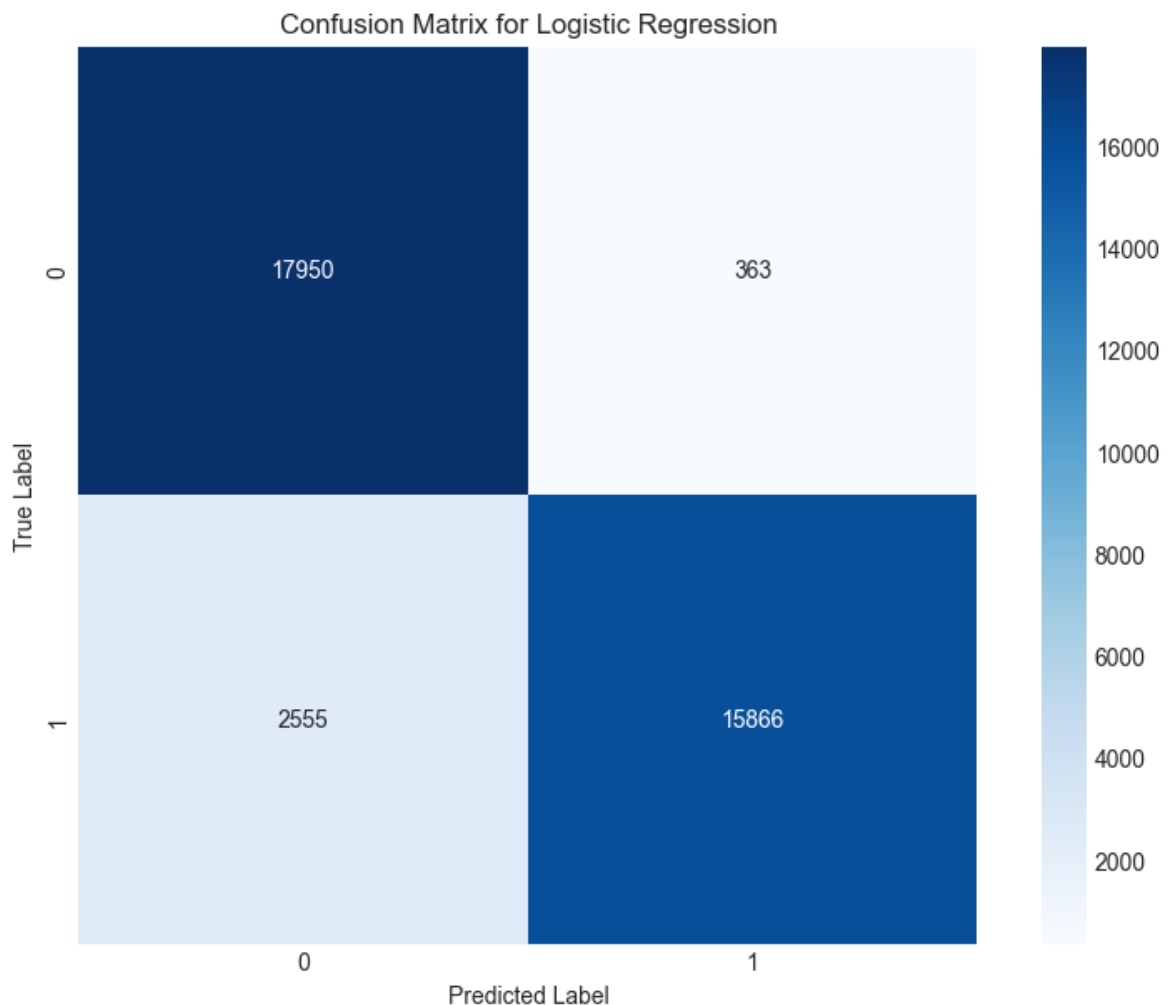
```
Out [38]: LogisticRegression
LogisticRegression(random_state=809)
```


Confusion Matrix

```
In [39]: # Predictions
y_pred = logit.predict(X_test)

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# Visualization of the Confusion Matrix
plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', square=True)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix for Logistic Regression')
plt.show()
```



ROC Curve & Precision-Recall


```
In [40]: # ROC Curve
y_pred_proba = logit.predict_proba(X_test)[:,1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = roc_auc_score(y_test, y_pred_proba)

plt.figure(figsize=(10, 5))
plt.plot(fpr, tpr, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Logistic Regression')
plt.legend(loc='best')
plt.show()

# Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_pred_proba)

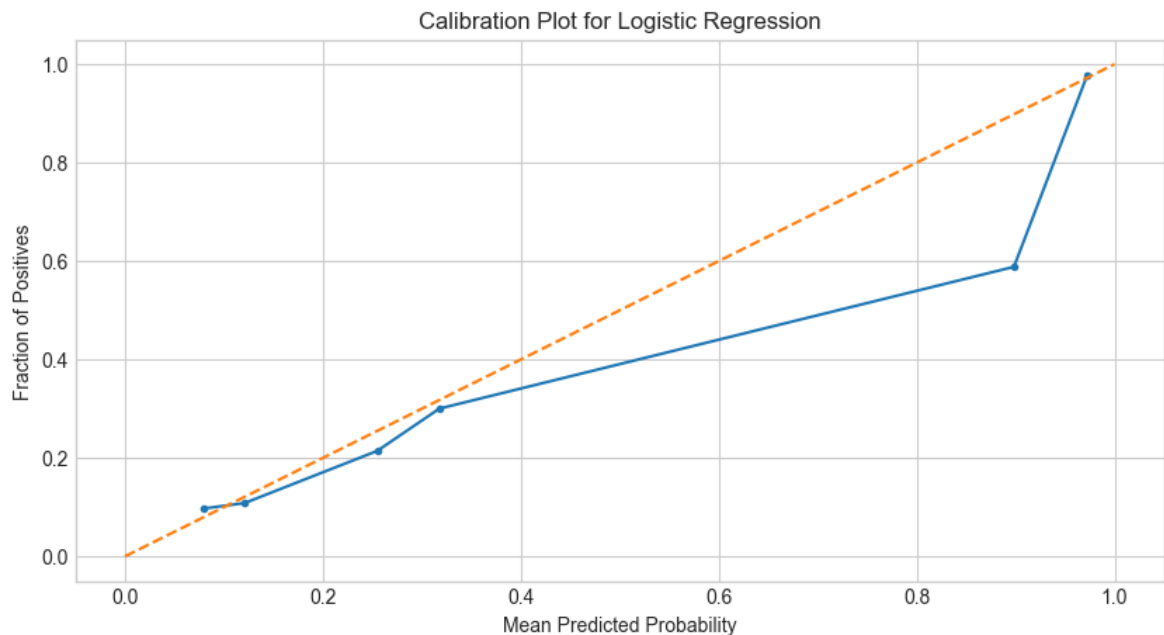
plt.figure(figsize=(10, 5))
plt.plot(recall, precision, marker='.')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve for Logistic Regression')
plt.show()
```



Calibration Plot

```
In [41]: # Calibration Plot
prob_true, prob_pred = calibration_curve(y_test, y_pred_proba, n_bins=10)
```

```
plt.figure(figsize=(10, 5))
plt.plot(prob_pred, prob_true, marker='.')
plt.plot([0, 1], [0, 1], linestyle='--')
plt.xlabel('Mean Predicted Probability')
plt.ylabel('Fraction of Positives')
plt.title('Calibration Plot for Logistic Regression')
plt.show()
```



KNN

Preparation

```
In [42]: # Initialize the KNN classifier, choosing an appropriate number of neighbors
knn = KNeighborsClassifier(n_neighbors=max_k)

# Fit the model
knn.fit(X_train, y_train)
```

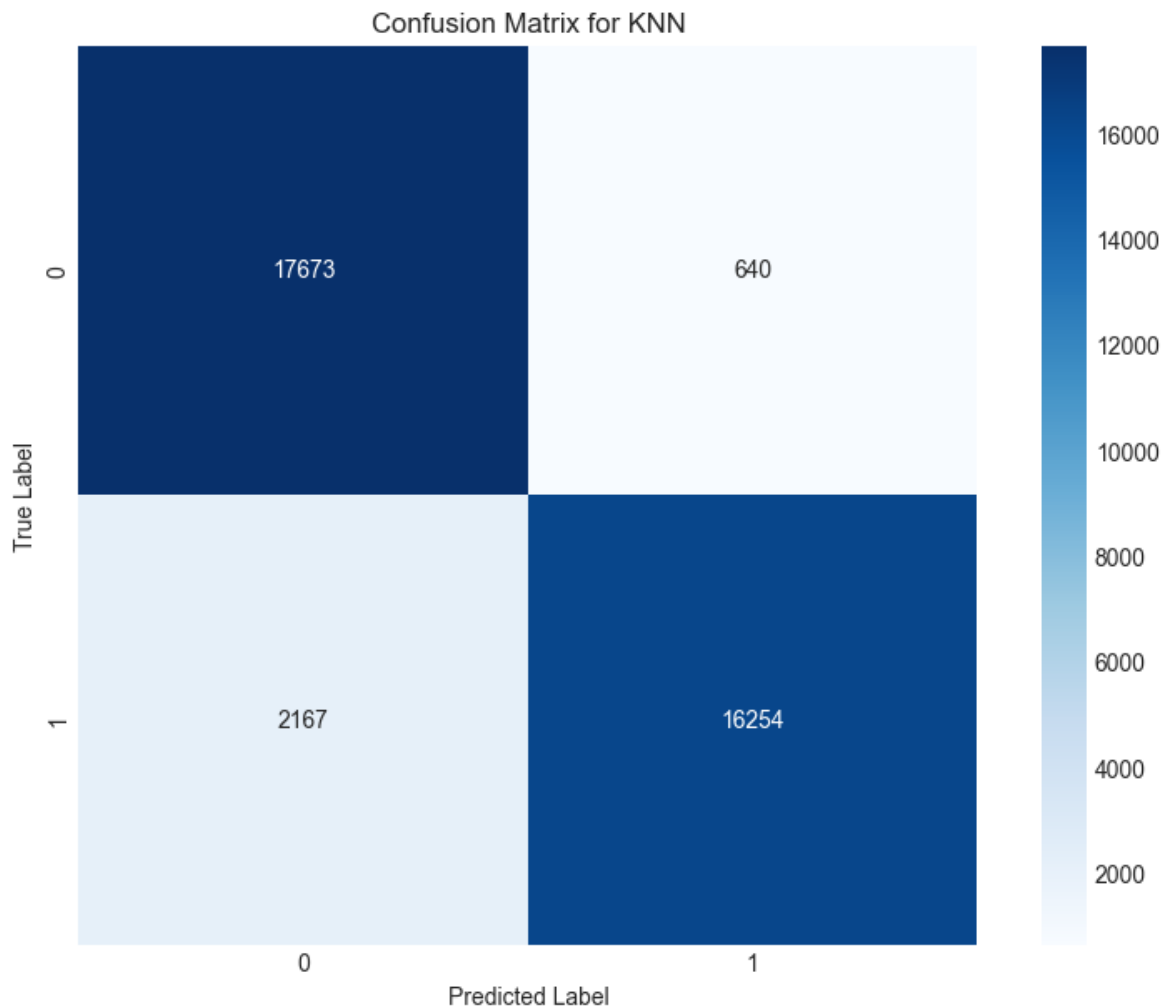
```
Out [42]: KNeighborsClassifier
KNeighborsClassifier(n_neighbors=4)
```

Confusion Matrix

```
In [43]: # Predictions
y_pred = knn.predict(X_test)

# Confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# Visualization
plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', square=True)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix for KNN')
plt.show()
```



ROC Curve & Precision-Recall


```
In [44]: # ROC Curve
y_pred_proba = knn.predict_proba(X_test)[:,-1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = roc_auc_score(y_test, y_pred_proba)

plt.figure(figsize=(10, 5))
plt.plot(fpr, tpr, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for KNN')
plt.legend(loc='best')

# Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_pred_proba)

plt.figure(figsize=(10, 5))
plt.plot(recall, precision, marker='.')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve for KNN')
```

Out[44]: Text(0.5, 1.0, 'Precision-Recall Curve for KNN')

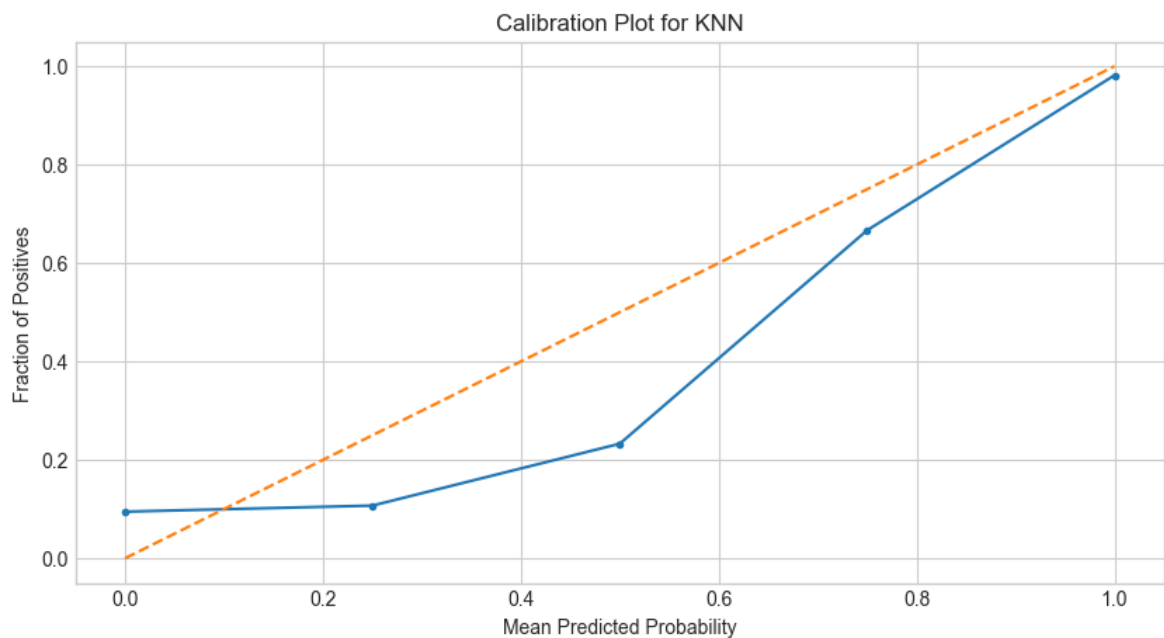


Calibration plot

```
In [45]: # Calibration plot
prob_true, prob_pred = calibration_curve(y_test, y_pred_proba, n_bins=10)

plt.figure(figsize=(10, 5))
plt.plot(prob_pred, prob_true, marker='.')
plt.plot([0, 1], [0, 1], linestyle='--')
plt.xlabel('Mean Predicted Probability')
plt.ylabel('Fraction of Positives')
plt.title('Calibration Plot for KNN')
```

```
Out[45]: Text(0.5, 1.0, 'Calibration Plot for KNN')
```



Random Forest

Preparation

```
In [46]: # Initialize the Random Forest classifier
rf = RandomForestClassifier(n_estimators=100, random_state=random_state)

# Fit the model
rf.fit(X_train, y_train)
```

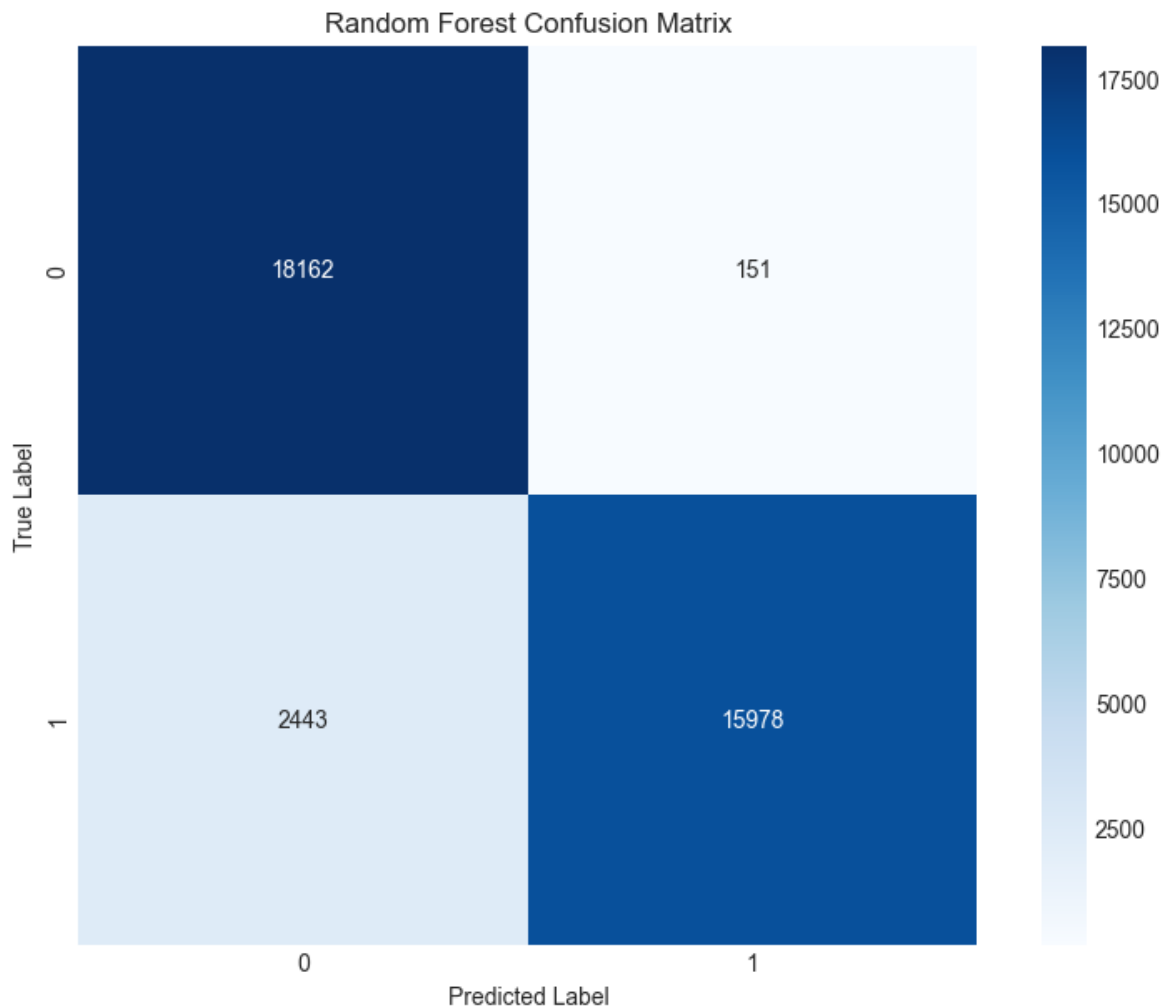
```
Out[46]: ▼ RandomForestClassifier
RandomForestClassifier(random_state=809)
```

Confusion Matrix

```
In [47]: # Predictions
y_pred = rf.predict(X_test)

# Confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# Visualization
plt.figure(figsize=(10, 7)) # Adjusts the figure size for better readability
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', square=True)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Random Forest Confusion Matrix')
plt.show()
```



ROC Curve & Precision-Recall


```
In [48]: # ROC Curve
y_pred_proba = rf.predict_proba(X_test)[:,-1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = roc_auc_score(y_test, y_pred_proba)

plt.figure(figsize=(10, 5))
plt.plot(fpr, tpr, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='best')

# Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_pred_proba)

plt.figure(figsize=(10, 5))
plt.plot(recall, precision, marker='.')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
```

```
Out[48]: Text(0.5, 1.0, 'Precision-Recall Curve')
```




Calibration Plot

```
In [49]: # Calibration plot
prob_true, prob_pred = calibration_curve(y_test, y_pred_proba, n_bins=10)

plt.figure(figsize=(10, 5))
plt.plot(prob_pred, prob_true, marker='.')
plt.plot([0, 1], [0, 1], linestyle='--')
plt.xlabel('Mean Predicted Probability')
plt.ylabel('Fraction of Positives')
plt.title('Calibration Plot')
```

```
Out[49]: Text(0.5, 1.0, 'Calibration Plot')
```

Xtreme Gradient Boosting (XGBoost)

Preparation

```
In [50]: # Initialize XGBoost Classifier
xgb_model = XGBClassifier(use_label_encoder=False, eval_metric='logloss',

# Fit the model
xgb_model.fit(X_train, y_train)
```

```
Out[50]: XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_
              rounds=None,
              enable_categorical=False, eval_metric='logloss',
              feature_types=None, gamma=None, grow_policy=None,
              importance_type=None, interaction_constraints=None,
              learning_rate=None, max_bin=None, max_cat_threshold
              =None,
```

Confusion Matrix

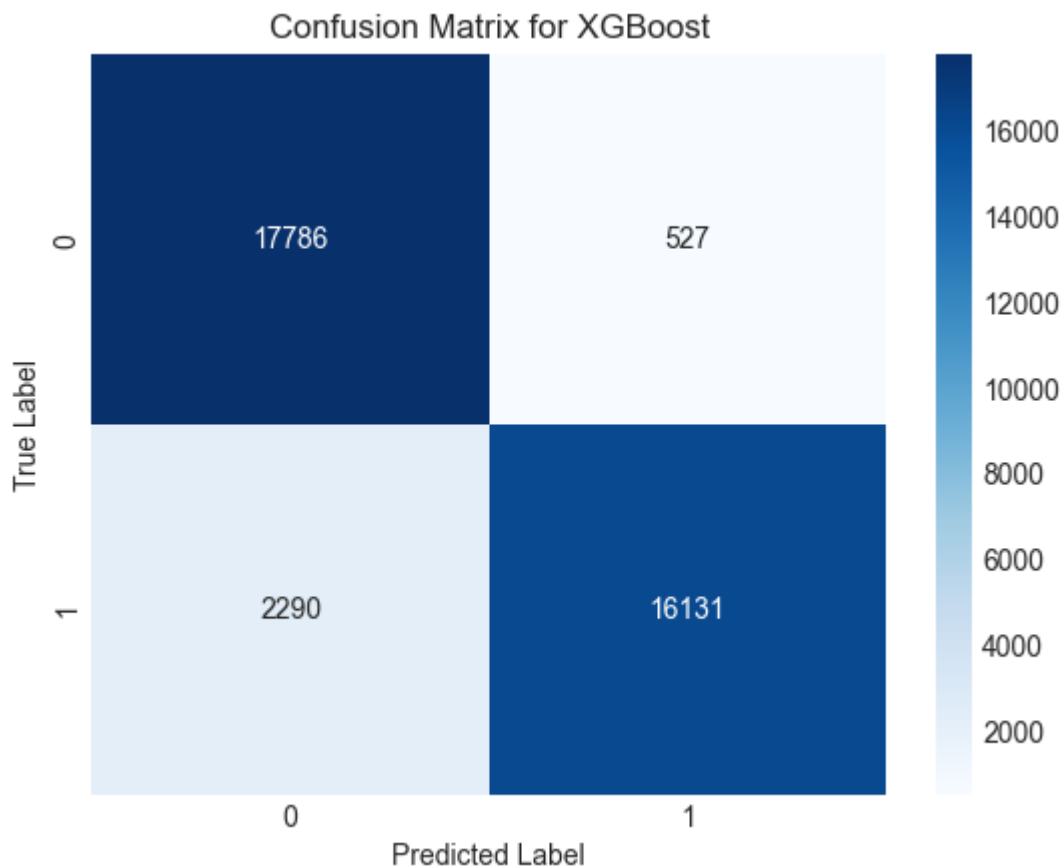
```
In [51]: # Predictions
y_pred = xgb_model.predict(X_test)

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# Visualization
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues")
plt.xlabel('Predicted Label')
```

```
plt.ylabel('True Label')
plt.title('Confusion Matrix for XGBoost')
```

Out[51]: Text(0.5, 1.0, 'Confusion Matrix for XGBoost')



ROC Curve & Precision-Recall


```
In [52]: # ROC Curve
y_pred_proba = xgb_model.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(10, 5))
plt.plot(fpr, tpr, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc="lower right")

# Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_pred_proba)

plt.figure(figsize=(10, 5))
plt.plot(recall, precision, marker='.')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
```

Out[52]: Text(0.5, 1.0, 'Precision-Recall Curve')



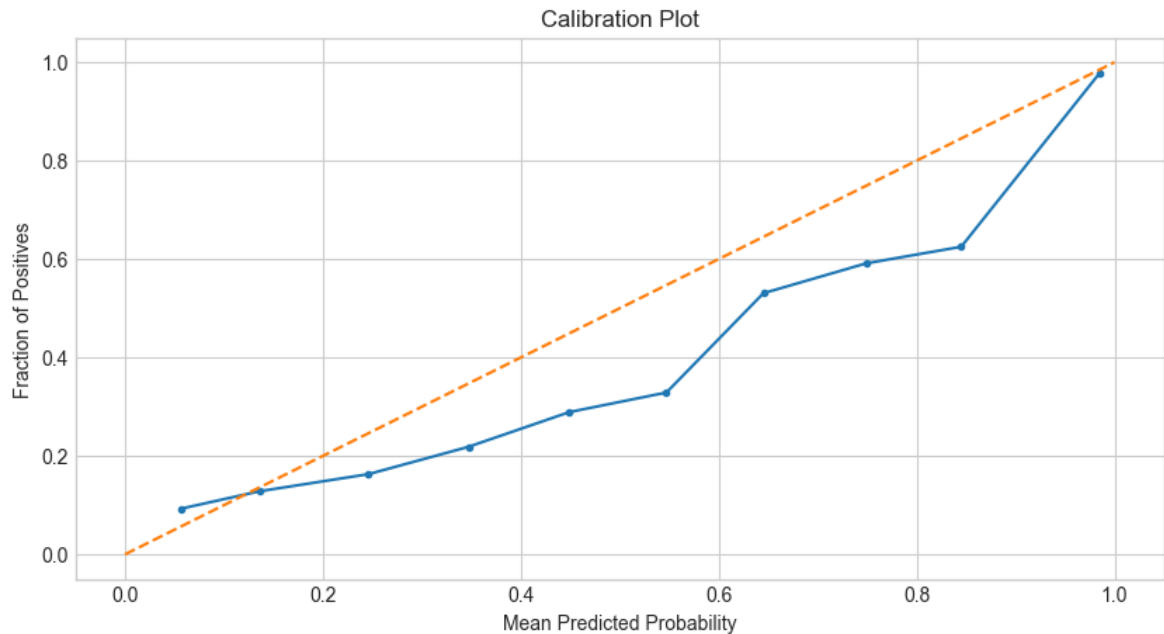
Calibration Plot

```
In [53]: # Calibrated probabilities
calibrated = CalibratedClassifierCV(xgb_model, method='sigmoid', cv='pref')
calibrated.fit(X_train, y_train)
prob_pos = calibrated.predict_proba(X_test)[:, 1]

# Calibration curve
prob_true, prob_pred = calibration_curve(y_test, prob_pos, n_bins=10)

plt.figure(figsize=(10, 5))
plt.plot(prob_pred, prob_true, marker='.')
plt.plot([0, 1], [0, 1], linestyle='--')
plt.xlabel('Mean Predicted Probability')
plt.ylabel('Fraction of Positives')
plt.title('Calibration Plot')
```

Out[53]: Text(0.5, 1.0, 'Calibration Plot')



Compare ROC curve

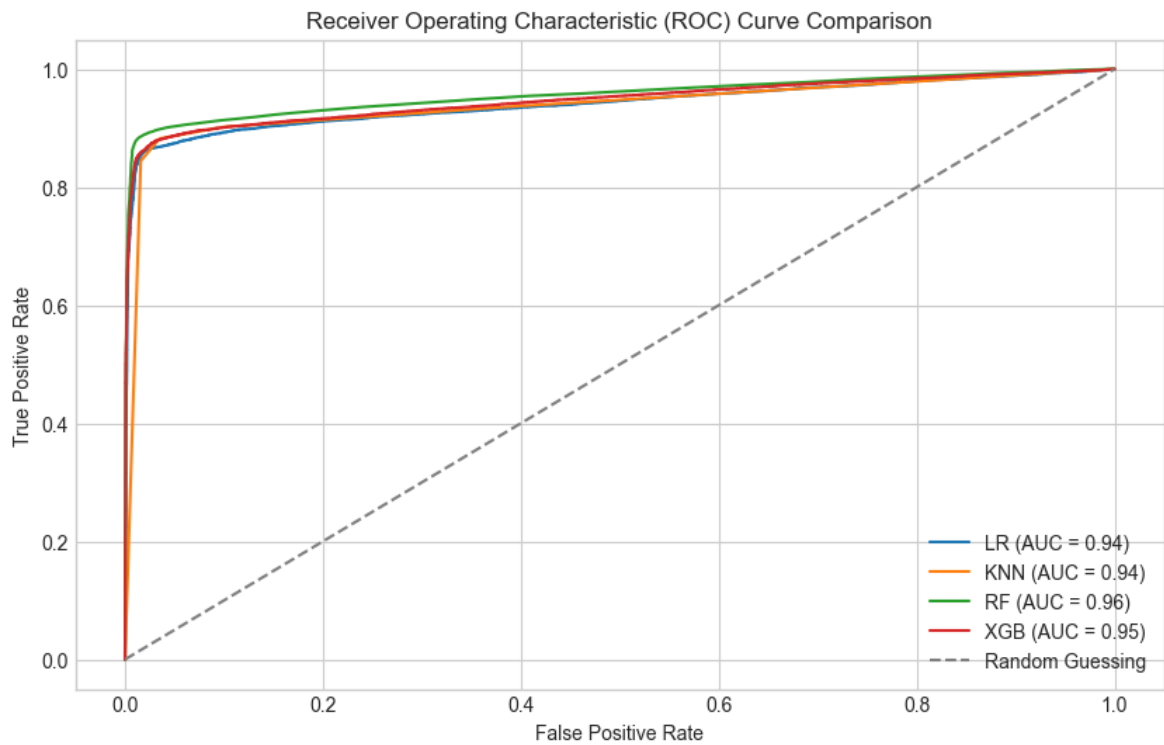
```
In [54]: # Initialize models for ROC curve comparison
models_roc = [('LR', LogisticRegression(random_state=random_state)),
               ('KNN', KNeighborsClassifier(n_neighbors=max_k)),
               ('RF', RandomForestClassifier(n_estimators=100, random_state=random_state)),
               ('XGB', XGBClassifier(use_label_encoder=False, eval_metric='logit'))]

# Plot ROC curve for each model
plt.figure(figsize=(10, 6))
for name, model in models_roc:
    model.fit(X_train, y_train) # Ensure models are trained before calculation
    y_pred_proba = model.predict_proba(X_test)[:, 1]
    fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, label=f'{name} (AUC = {roc_auc:.2f})')

# Plot the random guessing line
plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random Guess')

# Customize the plot
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve Comparison')
plt.legend(loc='lower right')

# Show the plot
plt.show()
```



Compare Precision-Recall curve

```
In [55]: # Initialize models for Precision-Recall curve comparison
models_pr = [('Logistic Regression', logit), ('KNN', knn), ('Random Fores

# Plot Precision-Recall curve for each model
plt.figure(figsize=(10, 6))
for name, model in models_pr:
    model.fit(X_train, y_train) # Ensure models are trained before calcula
    y_pred_proba = model.predict_proba(X_test)[: , 1]
    precision, recall, _ = precision_recall_curve(y_test, y_pred_proba)
    pr_auc = auc(recall, precision)
    plt.plot(recall, precision, label=f'{name} (AUC = {pr_auc:.2f})')

# Customize the plot
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve Comparison')
plt.legend()

# Show the plot
plt.show()
```

